

CSE 421/521 - Operating Systems
Fall 2019

RECITATION - II
INTRODUCTION TO PINTOS
PROF. TEVFIK KOSAR

Presented by Naveena Alango

University at Buffalo
September 16-19th, 2019

Verify Setup

- Compile
 - `$ cd $PINTOSDIR/src/threads`
 - `$ make`
- Test
 - `$ cd build`
 - `$ pintos run alarm-multiple`
- New window popped up?

Get Familiar with the Code

- The first task is to read and understand the code for the initial thread system (under the “[pintos/src/threads/](#)” directory).
- Pintos already implements thread creation and thread completion, a simple scheduler to switch between threads, and synchronization primitives (semaphores, locks, condition variables, and optimization barriers).
- For a brief overview of the files in the “threads/” directory, please see “[Section 2.1.2 Source Files](#)” in the Pintos Reference Guide

Pintos Thread System

- Read threads/thread.c and threads/synch.c to understand
 - How the switching between threads occur
 - How the provided scheduler works
 - How the various synchronizations primitives work


Pintos Thread System

Defined in threads/thread.h:

```
struct thread
{
    tid_t tid;          /* Thread identifier. */
    enum thread_status status; /* Thread state. */
    char name[16]; /* Name (for debugging purposes). */
    uint8_t *stack; /* Saved stack pointer. */
    int priority; /* Priority. */
    struct list_elem allelem; /* List element for all-threads list. */
    /* Shared between thread.c and synch.c. */
    struct list_elem elem; /* List element. */
};
```

You add more fields here as you need them.

```
#ifndef USERPROG
/* Owned by userprog/process.c. */
uint32_t *pagedir; /* Page directory. */
#endif
/* Owned by thread.c. */
unsigned magic; /* Detects stack overflow. */
};
```



Task 1: Implement Alarm Clock

- Reimplement `timer_sleep()` in `devices/timer.c` without busy waiting

```
/* Suspends execution for approximately TICKS timer ticks. */
```

```
void timer_sleep (int64_t ticks){  
    int64_t start = timer_ticks ();  
    ASSERT (intr_get_level () == INTR_ON);  
    while (timer_elapsed (start) < ticks)  
        thread_yield ();  
}
```

- Implementation details
 - Remove thread from ready list and put it back after sufficient ticks have elapsed

=> Get rid of the while loop!

=> Put the thread in to the waiting queue

=> `sema_down()` // the thread will be waken up when `sema_up()` is called

Any implementation using busy-waiting will not get full credits!

Task 2A: Implement Priority Scheduler

- Ready thread with highest priority gets the processor
- When a thread is added to the ready list that has a higher priority than the currently running thread, immediately yield the processor to the new thread
- When threads are waiting for a lock, semaphore or a condition variable, the highest priority waiting thread should be woken up first
- Implementation details
 - compare priority of the thread being added to the ready list with that of the running thread
 - select next thread to run based on priorities
 - compare priorities of waiting threads when releasing locks, semaphores, condition variables

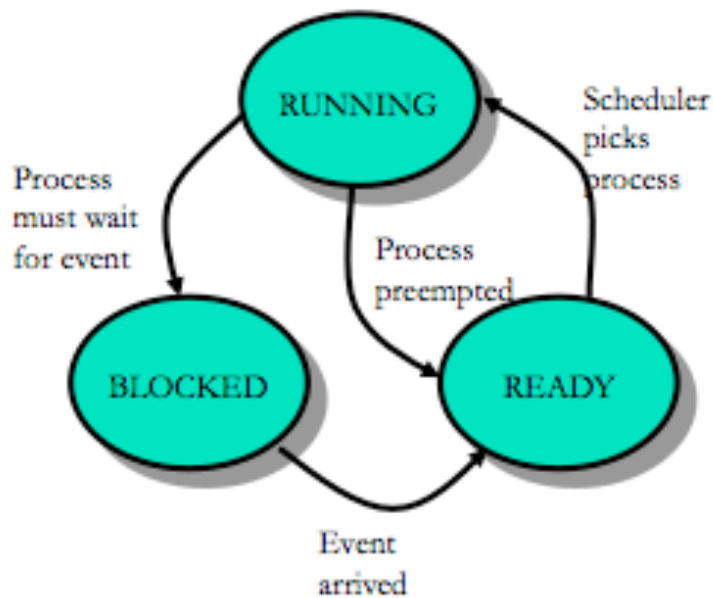
 **Preemptive**

=> The default scheduler is very primitive, FCFS, and needs to be changed

=> Search the threads according to their priorities, and release highest...

=> You can reuse code from lib/kernel/list.c

use `thread_yield()` to implement preemption



- Current thread (“**RUNNING**”) is moved to **READY** state, added to **READY** list.
- Then scheduler is invoked. Picks a new **READY** thread from **READY** list.
- Case a): there’s only 1 **READY** thread. Thread is rescheduled right away
- Case b): there are other **READY** thread(s)
 - b.1) another thread has higher priority – it is scheduled
 - b.2) another thread has same priority – it is scheduled provided the previously running thread was inserted in tail of ready list.
- “`thread_yield()`” is a call you can use whenever you identify a need to preempt current thread.
- **Exception:** inside an interrupt handler, use “`intr_yield_on_return()`” instead

* re-implement **`next_thread_to_run()`** for priority scheduling

Priority Inversion

- Strict priority scheduling can lead to a phenomenon called “priority inversion”
- Supplemental reading:
 - What really happened to the Pathfinder on Mars?
- Consider the following example where $\text{prio}(H) > \text{prio}(M) > \text{prio}(L)$
 - H needs a lock currently held by L, so H blocks
 - M that was already on the ready list gets the processor before L
 - H indirectly waits for M
 - (on Path Finder, a watchdog timer noticed that H failed to run for some time, and continuously reset the system)

Task 2B: Implement Priority Donation

- When a high priority thread H waits on a lock held by a lower priority thread L, donate H's priority to L and recall the donation once L releases the lock
- Implement priority donation for locks
- Important:
 - Remember to return L to previous priority once it releases the lock.
 - Be sure to handle multiple donations (max of all donations)
 - Be sure to handle nested donations, e.g., H waits on M which waits on L...

Synchronization

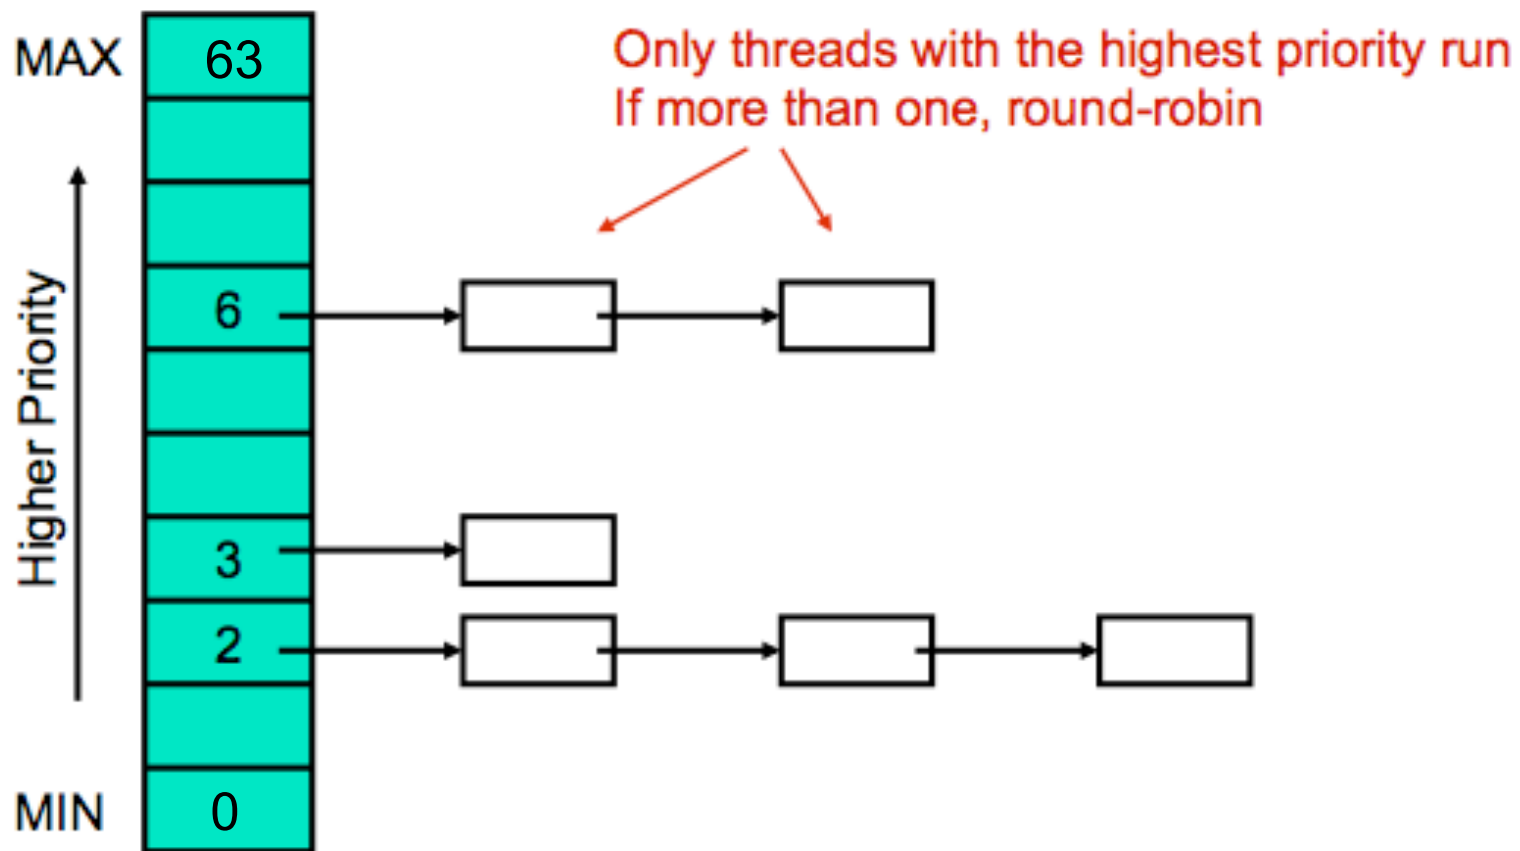
- Any synchronization problem can be easily solved by turning interrupts off: while interrupts are off, there is no concurrency, so there's no possibility for race conditions. **But, you should NOT do this!**
- Instead, **use semaphores, locks, and condition variables** to solve the bulk of your synchronization problems.
- **Any implementation turning the interrupts off for synchronization purposes, will not get full credits!**
- **The only place you are allowed to turn interrupts off is,** when coordinating data shared between a kernel thread and an interrupt handler. Because interrupt handlers can't sleep, they can't acquire locks.

Task 3: Implement Advanced Scheduler

- Implement Multi Level Feedback Queue Scheduler
- Priority donation not needed in the advanced scheduler – two implementations are not required to coexist
 - Only one is active at a time
- Advanced Scheduler must be chosen only if ‘`–mlfq`’ kernel option is specified
- Read section on 4.4 BSD Scheduler in the Pintos manual for detailed information
- Some of the parameters are real numbers and calculations involving them have to be simulated using integers.
 - Write a fixed-point layer (header file)

4.4BSD Priority Based Scheduling

4.4BSD scheduler has 64 priorities and thus 64 ready queues, numbered 0 (PRI_MIN) through 63 (PRI_MAX).



Calculating Priority

- NOTE: Lower numbers correspond to lower priorities in 4.4BSD, so that priority 0 is the lowest priority and priority 63 is the highest.
- Every 4 clock ticks, calculate:
$$\text{priority} = \text{PRI_MAX} - (\text{recent_cpu} / 4) - (\text{nice} * 2)$$

(rounded down to the nearest integer)
- It gives a thread that has received CPU time recently lower priority for being reassigned the CPU the next time the scheduler runs.

← Aging

Nice Value

==> how “nice” the thread should be to other threads.

- A nice of zero does not affect thread priority.
- A positive nice, to the **maximum of 20**, decreases the priority of a thread and causes it to give up some CPU time it would otherwise receive.

A negative nice, to the **minimum of -20**, tends to take away CPU time from other threads.

Calculating recent_cpu

- will be implemented in `thread_get_recent_cpu(void)`
- An array of n elements to track the CPU time received in each of the last n seconds requires $O(n)$ space per thread and $O(n)$ time per calculation of a new weighted average.
- Instead, we use an exponentially weighted moving average:
 - **`recent_cpu(0) = 0`** *// or parent thread's value*
 - *at each timer interrupt, **`recent_cpu++`** for the running thread.*
 - *and once per second, for each thread:*

$$\text{recent_cpu}(t) = a * \text{recent_cpu}(t-1) + \text{nice}$$

$$\text{where, } a = (2 * \text{load_avg}) / (2 * \text{load_avg} + 1)$$

Calculating load_avg

- will be implemented in `thread_get_load_avg(void)`
- Estimates the average number of threads ready to run over the past minute.
- Like `recent_cpu`, it is an exponentially weighted moving average.
- Unlike `priority` and `recent_cpu`, `load_avg` is system-wide, not thread-specific.
- At system boot, it is **initialized to 0**. Once per second thereafter, it is updated according to the following formula:

$$\text{load_avg}(t) = (59/60) * \text{load_avg}(t-1) + (1/60) * \text{ready_threads}$$

- `ready_threads`: number of threads that are either running or ready to run at the time of update

Functions to implement

- Skeletons of these functions are provided in “threads/threads.c”:
- `int thread_get_nice (void)`
- `void thread_set_nice (int new_nice)`
- `void thread_set_priority (int new_priority)`
- `int thread_get_priority (void)`
- `int thread_get_recent_cpu (void)`
- `int thread_get_load_avg (void)`

Suggested Order of Implementation

- Alarm Clock
 - easier to implement compared to the other parts
 - other parts not dependent on this
- Priority Scheduler
 - needed for implementing Priority Donation and Advanced Scheduler
- Priority Donation | Advanced Scheduler
 - these two parts are independent of each other
 - can be implemented in any order but only after Priority Scheduler is ready

Debugging

Find a line in a file recursively using **grep**

```
ex) $ grep -rin "timer_sleep"
```

- **ASSERT**
 - *Defined in*
 - *ASSERT(expression): Test the value of expression. If it evaluates to zero (false), the kernel panics.*
- Using **printf**
 - You can call it from (practically) anywhere in the kernel
- Using **GDB**

Debugging with GDB

- On first terminal:

```
$ pintos --gdb -- run alarm-multiple
```

- In another terminal

```
$ pintos-gdb kernel.o
```

- (gdb) target remote localhost:1234 OR
- (gdb) debugpintos

GDB

- **Continue**
 - (gdb) c
- **Break**
 - (gdb) break function
 - (gdb) break file:line
- **Print**
 - (gdb) p expression
- **Next**
 - (gdb) n
- **List**
 - (gdb) l

Testing

- Make grade

- Possible to see actual pintos running command
- `./threads/build/grade` // contains summary of all tests.

- `threads/build/tests/threads/`

- `*.error`
- `*.output`
- `*.result`

Expected Lines of Code

devices/timer.c		42	+++++-
threads/fixed-point.h		120	+++++
threads/synch.c		88	+++++++-
threads/thread.c		196	+++++-----
threads/thread.h		23	+++

- 5 files changed, 440 insertions(+), 29 deletions(-)
- ***'fixed-point.h'*** is a new file added by the reference solution.